

## Assignment 7

### Object-Relational Database Programming

#### 1. In this problem, you can not use arrays.

Consider the relation schema `Graph(source int, target int)` representing the schema for storing a directed graph  $G$  as a set of edges.

Recall that  $G$  is *connected* if for each pair of vertices  $(s, t)$  in  $G$  with  $s \neq t$ , there exists a path in  $G$  from  $s$  to  $t$ , i.e., if for each pair of vertices  $(s, t)$  in  $G$  with  $s \neq t$ ,  $(s, t)$  is in the transitive closure of  $G$ .

An *articulation vertex* of  $G$  is a vertex  $\mathbf{v}$  of  $G$  such that removing the edges in  $G$  with source or target  $\mathbf{v}$  results in a graph that is **not** connected. More formally,  $\mathbf{v}$  is an articulation vertex of  $G$ , if the graph

$$G - (\{(s, t) | (s, t) \in G \wedge (s = \mathbf{v} \vee t = \mathbf{v})\})$$

is **not** connected.

We say that  $G$  is *bi-connected* if  $G$  does not have any articulation vertices.

Write a PostgreSQL program `biConnected()` that returns true if  $G$  is bi-connected and false otherwise.

#### 2. In this problem, you can not use arrays.

Implement the HITS authority-hubs algorithm which is a variant of a page ranking algorithm. For more information about the HITS algorithm consult

[https://en.wikipedia.org/wiki/HITS\\_algorithm](https://en.wikipedia.org/wiki/HITS_algorithm)  
[https://www.youtube.com/watch?v=jr3YGgfDY\\_E](https://www.youtube.com/watch?v=jr3YGgfDY_E)

The input data is given in a relation `Graph(source integer, target integer)` which represent the graph on which the HITS Algorithm operates.<sup>1</sup> So each vertex in this graph will receive an authority and a hub score.

An important detail of the HITS algorithm concerns the normalization of the authority vector (analogously, the hub vector). This vector needs to be normalized to have norm = 1 after each iteration step. Otherwise, the algorithm will not converge.

Normalization of a vector of numbers can be done as follows: If  $\mathbf{x} = (x_1, \dots, x_n)$  is a vector of real numbers, then its norm  $|\mathbf{x}|$  is given by the formula  $\sqrt{x_1^2 + \dots + x_n^2}$ . Therefore, you can normalize the vector  $(x_1, \dots, x_n)$  by transforming it to the vector  $\frac{\mathbf{x}}{|\mathbf{x}|} = (\frac{x_1}{|\mathbf{x}|} + \dots + \frac{x_n}{|\mathbf{x}|})$ . The norm of this vector will be 1.

---

<sup>1</sup>This graph is called the *base* in the Wikipedia article.

3. **In this problem, you can use arrays, but only as a mechanism to represents sets of words.**

Consider the relation schema `document(doc int, words text[])` representing a relation of pairs  $(d, W)$  where  $d$  is a unique document id and  $W$  denotes the set of words that occur in  $d$ .

Let  $\mathbf{W}$  denote the set of all words that occur in the documents and let  $t$  be a positive integer denoting a *threshold*.

Let  $X \subseteq \mathbf{W}$ . We say that  $X$  is  $t$ -frequent if

$$\text{count}(\{d \mid (d, W) \in \text{document and } X \subseteq W\}) \geq t$$

In other words,  $X$  is  $t$ -frequent if there are at least  $t$  documents that contain all the words in  $X$ .

Write a PostgreSQL program `frequentSets(t int)` that returns the relation of all  $t$ -frequent sets.

In a good solution for this problem, you should use the following rule: if  $X$  is not  $t$ -frequent then any set  $Y$  such that  $X \subseteq Y$  is not  $t$ -frequent either. In the literature, this is called the *Apriori* rule of the frequent itemset mining problem. This rule allows you to avoid examining supersets of sets that are not frequent. This can drastically reduce the search space.

4. **In this problem you can not use arrays.**

Suppose you have a weighted undirected graph  $G$  stored in a ternary table with schema

`Graph(source int, target int, weight int)`

A triple  $(s, t, w)$  in `Graph` indicates that *Graph* has an edge  $(s, t)$  whose edge weight is  $w$ . (In this problem, we will assume that each edge weight is a positive integer.)

Since the graph is undirected, whenever there is an weighted edge  $(s, t, w)$  in  $G$ , then  $(t, s, w)$  is also a weighted edge in the  $G$ . Below is an example of a graph  $G$ .

Graph  $G$

source	target	weight
0	1	2
1	0	2
0	4	10
4	0	10
1	3	3
3	1	3
1	4	7
4	1	7
2	3	4
3	2	4
3	4	5
4	3	5
4	2	6
2	4	6

Implement Dijkstra's Algorithm as a PostgreSQL function `Dijkstra(s integer)` to compute the shortest path lengths (i.e., the distances) from some input vertex  $s$  in  $G$  to all other vertices in  $G$ . `Dijkstra(s integer)` should accept an argument  $s$ , the source vertex, and outputs a table `Paths` which represents the pairs  $(t, d)$  where  $d$  is the shortest distance from  $s$  to  $t$ . To test your procedure, you can use the graph shown above.

You can find a description of Dijkstra's algorithm as the following webpage [https://en.wikipedia.org/wiki/Dijkstra%27s\\_algorithm#Pseudocode](https://en.wikipedia.org/wiki/Dijkstra%27s_algorithm#Pseudocode)

When you apply `Dijkstra(0)`, you should obtain the following `Paths` table:

target	distanceToTarget
0	0
1	2
2	9
3	5
4	9

5. **In this problem, you can not use arrays.**

Consider the relation schema `Graph(source int, target int)` representing the schema for storing a directed graph  $G$  of edges.

Let 'red', 'green', and 'blue' be 3 colors. We say that  $G$  is *3-colorable* if it is possible to assign to each vertex of  $G$  one of these 3 colors provided that, for each edge  $(s, t)$  in  $G$ , the color assigned to  $s$  is different than the color assigned to  $t$ .

Write a PostgreSQL program `threeColorable()` that returns true if  $G$  is 3-colorable, and false otherwise.

(Hint: use a backtracking algorithm that finds a 3-color assignment to the vertices of  $G$  if such an assignment exists.)